
Instacart Case Study



Online Groceries

Overview

Purpose & Context

Whilst studying with Career Foundry I completed a project that introduced me to Python and using it to perform an exploratory analysis in order to derive insights and suggest strategies. It focused on Instacart, an online grocery service.



Project Goal

Uncover more information about sales patterns. Perform an initial data and exploratory analysis in order to derive insights and suggest strategies for better segmentation based on the provided criteria

Key Objectives

Identify busiest days of the week and hours of the day for ad scheduling
Establish times of the day when people spend the most money
Which departments have the highest frequency of product orders
Uncover customer ordering behaviours based on various features such as region, family status, loyalty, and other demographics

Overview

Datasets

Instacart Data Sets:

- Data Dictionary
- "The Instacart Online Grocery Shopping Dataset 2017", Accessed from www.instacart.com/datasets/grocery-shopping-2017 via Kaggle on 23.6.23

CareerFoundry Data Sets:

- Customers Data Set

Skills

Python

Data wrangling

Data merging

Deriving variables

Grouping data

Aggregating data

Reporting in Excel

Population flows

Tools



python™



Steps





Data Preparation

Overview

- Basic descriptive exploratory tasks
- Changing data types
- Access values and determine their meaning using a data dictionary

- Download data and import into notebook as a pandas dataframe
- Conduct basic descriptive exploratory tasks
- Change data types of identifier variables into more suitable types and rename columns where needed
- Access values and determine their meaning using a data dictionary

The first step was to clean the 3 datasets (orders, products, customers) if you skip this stage it can result in misleading results. It involved checking for missing values, inconsistent data types, and removing unnecessary columns. Python proved to be a great, and fast, tool for this process.



Analysis Process

Overview

- Creating new dataframes
- Data cleaning
- Merging dataframes
- Creating new columns
- Creating flags
- Creating summary columns
- Creating visualisations

- Creating new dataframes based on a certain criteria
- Fixing mixed-type variables, missing values, and removing duplicates
- Merging sets of dataframes and exporting as pickle
- Creating new columns using conditional logic in the form of if-statements, user-defined functions, the loc() function, and for-loops
- Creating flags
- Creating summary columns of descriptive statistics using the groupby() function
- Creating histograms, bar charts, line charts, and scatterplots
- Putting together a report

Merging the dataframes proved to be the most challenging aspect of this process as I was dealing with 32 million rows of data. I encountered a few issues with freezing due to processing power when working on the dataframe, so eventually had to split the data into a training and test set (70/30). I also changed the data types to reduce the size of the dataframe, making it as streamlined as possible. Moving forward I used this smaller set to perform a lot of the analysis. This also prevented crashing issues when creating visualisations. In the future I would have to use a more powerful machine.



Results & Recommendations

Overview

- Schedule advertising
- Grouping products by mark-up
- Promotions
- Incentivise brand loyalty.
- Targeted advertising.

From the analysis process I was able to come up the following recommendations,


- Focus advertising high-end products during early hours targeting commuters and high-earners
 - Groupings are based just on purchase price. They could also group products by mark-up and attribute values to them that way.
 - Attach promotions to these best selling items as they seem to be staples, which then encourage customers to buy less well selling items.
 - Incentivise brand loyalty by introducing repeat customer bonus schemes, a points system for example.
 - Those over 40 have more spending power so targeting advertising to this age group could help boost income.
-



Visualisations


In the following pages are some of the visualisations

AGE PROFILE




This pie chart shows that the majority of customers are aged under 65.

PROFILE OF REGION AND LOYALTY




The most loyal customers are in the south.

PERCENTAGE OF REGION, GROUP BY DEPARTMENT



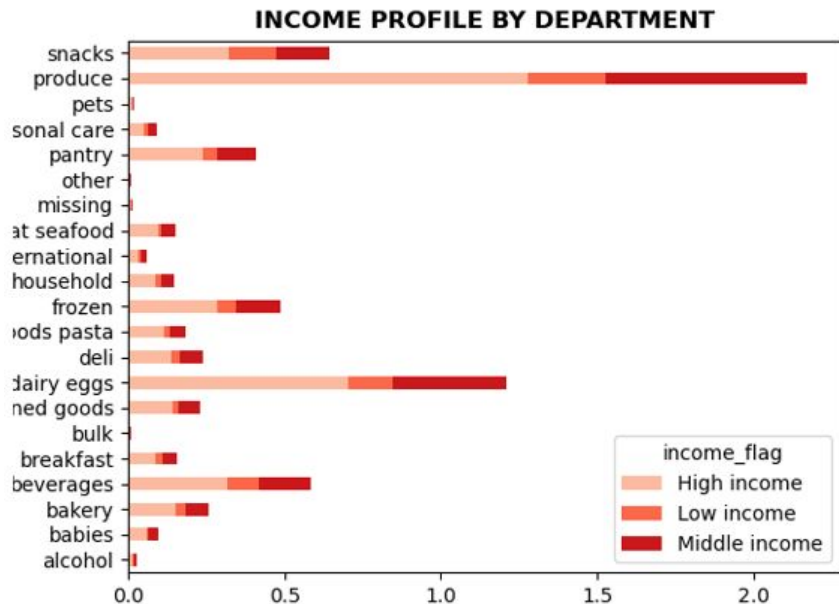
Scatterplot



This scatterplot shows the relationship between customer age and income. It shows that the majority of people earn under 200k. Very few 20-40 year olds earn more than 200k, with none earning more than 400k per year. People older than 40 start making over 200k, with some earning up to 600k. From this chart we can see that those over 40 have more spending power.




Stacked Bar Chart



This shows that high income customers spend more per department than those earning less by roughly the same proportion as the income profile share (58%). An exception is snacks, where they only make up 50%, suggesting that those on lower incomes buy more snacks.



Line Plot Chart



This line plot chart shows the average price of items ordered by order hour of the day.
It seems people spend more between 3am and 5am.





Python Code

Below are some examples of code used throughout the analysis. Changing data types, creating cross-tabs and merging dataframes.

```
dtypes: category(1), float64(1), int64(6)
memory usage: 2.4 GB

In [21]: # Change types for orders products combined data set

df_or_pr_co['product_id'] = df_or_pr_co['product_id'].astype('int32')
df_or_pr_co['reordered'] = df_or_pr_co['reordered'].astype('int8')
df_or_pr_co['add_to_cart_order'] = df_or_pr_co['add_to_cart_order'].astype('int32')
df_or_pr_co['order_id'] = df_or_pr_co['order_id'].astype('int32')
df_or_pr_co['user_id'] = df_or_pr_co['user_id'].astype('int32')
df_or_pr_co['order_number'] = df_or_pr_co['order_number'].astype('int8')
df_or_pr_co['orders_day_of_week'] = df_or_pr_co['orders_day_of_week'].astype('int8')
df_or_pr_co['order_hour_of_day'] = df_or_pr_co['order_hour_of_day'].astype('int8')
df_or_pr_co['days_since_last_order'] = df_or_pr_co['days_since_last_order'].astype('float16')

In [22]: df_or_pr_co . info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 32434489 entries, 0 to 32434488
Data columns (total 10 columns):
#   Column      Dtype
---  ---
0   order_id    int32
1   user_id     int32
2   order_number int8
3   orders_day_of_week int8
4   order_hour_of_day int8
5   days_since_last_order float16
6   product_id  int32
7   add_to_cart_order int32
8   reordered   int8
   _merge      category
dtypes: category(1), float16(1), int32(4), int8(4)
memory usage: 958.9 MB

In [23]: df_prods . info()
```

```
In [64]: # changing data type of department_id in df_dep
df_dep['department_id'] = df_dep['department_id'].astype('int8')

In [66]: df_dep . info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21 entries, 0 to 20
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   department_id 21 non-null      int8
1   department     21 non-null      object
dtypes: int8(1), object(1)
memory usage: 317.0+ bytes

In [67]: #merging department data frame with df4 dataframe
df_final = df4.merge(df_dep, on='department_id')

In [68]: # creating a crosstab between department and age
department_age = pd.crosstab(df_final['department'], df_final['age_flag'], dropna=False)

In [69]: #creating a stacked bar chart of age and department
dept_age = department_age.plot.barh(stacked=True, colors=sns.color_palette('Reds', 3))
plt.title('AGE PROFILE BY DEPARTMENT', fontweight='bold')
plt.show()
```

AGE PROFILE BY DEPARTMENT



```
average_price    float32
spending_flag     object
median_last_orders float16
order_frequency_flag object
dtype: object

In [29]: #combining the data sets on the 'user_id' column
df_combined = ord_pro_merge(df, on = 'user_id') #crashed my computer

In [30]: df_merged = ord_pro_merge(df, on = 'user_id')

In [31]: #after many attempts at merging the dataframes in different way I have finally managed to merge a 30X sample with the cus

In [32]: df_merged . shape

Out[32]: (9721098, 33)

In [33]: df_merged . head()

Out[33]:
order_id user_id order_number orders_day_of_week order_hour_of_day days_since_last_order product_id add_to_cart_order reordered
0 2539329      1           1           2           8           NaN           196           1
1 473747       1           3           3           12          21.0          196           1
2 2254736       1           4           4           7           29.0          196           1
```

Thanks



<https://github.com/danfradat/Instacartpython/tree/main>
